

8.2.2 EXPLAIN Output Format

The [EXPLAIN](#) statement provides information about the execution plan for a [SELECT](#) statement.

[EXPLAIN](#) returns a row of information for each table used in the [SELECT](#) statement. It lists the tables in the output in the order that MySQL would read them while processing the statement. MySQL resolves all joins using a nested-loop join method. This means that MySQL reads a row from the first table, and then finds a matching row in the second table, the third table, and so on. When all tables are processed, MySQL outputs the selected columns and backtracks through the table list until a table is found for which there are more matching rows. The next row is read from this table and the process continues with the next table.

When the [EXTENDED](#) keyword is used, [EXPLAIN](#) produces extra information that can be viewed by issuing a [SHOW WARNINGS](#) statement following the [EXPLAIN](#) statement. See [Section 8.2.3, “EXPLAIN EXTENDED Output Format”](#).

- [EXPLAIN Output Columns](#)
- [EXPLAIN Join Types](#)
- [EXPLAIN Extra Information](#)
- [EXPLAIN Output Interpretation](#)

EXPLAIN Output Columns

This section describes the output columns produced by [EXPLAIN](#). Later sections provide additional information about the [type](#) and [Extra](#) columns.

Each output row from [EXPLAIN](#) provides information about one table. Each row contains the values summarized in [Table 8.1, “EXPLAIN Output Columns”](#), and described in more detail following the table.

Table 8.1 EXPLAIN Output Columns

| Column | Meaning |
|-------------------------------|---------------------------------------|
| id | The SELECT identifier |
| select_type | The SELECT type |
| table | The table for the output row |
| type | The join type |
| possible_keys | The possible indexes to choose |
| key | The index actually chosen |
| key_len | The length of the chosen key |
| ref | The columns compared to the index |
| rows | Estimate of rows to be examined |
| Extra | Additional information |

Section Navigation [\[Toggle\]](#)

- [8.2 Obtaining Query Execution Plan Information](#)
- [8.2.1 Optimizing Queries with EXPLAIN](#)
- 8.2.2 EXPLAIN Output Format
- [8.2.3 EXPLAIN EXTENDED Output Format](#)
- [8.2.4 Estimating Query Performance](#)

- **id**

The [SELECT](#) identifier. This is the sequential number of the [SELECT](#) within the query. The value can be **NULL** if the row refers to the union result of other rows. In this case, the **table** column shows a value like **<unionM,N>** to indicate that the row refers to the union of the rows with **id** values of **M** and **N**.

- **select_type**

The type of [SELECT](#), which can be any of those shown in the following table.

| select_type Value | Meaning |
|-----------------------------|---|
| SIMPLE | Simple SELECT (not using UNION or subqueries) |
| PRIMARY | Outermost SELECT |
| UNION | Second or later SELECT statement in a UNION |
| DEPENDENT UNION | Second or later SELECT statement in a UNION , dependent on outer query |
| UNION RESULT | Result of a UNION . |
| SUBQUERY | First SELECT in subquery |
| DEPENDENT SUBQUERY | First SELECT in subquery, dependent on outer query |
| DERIVED | Derived table SELECT (subquery in FROM clause) |
| UNCACHEABLE SUBQUERY | A subquery for which the result cannot be cached and must be re-evaluated for each row of the outer query |

DEPENDENT typically signifies the use of a correlated subquery. See [Section 13.2.9.7. “Correlated Subqueries”](#).

DEPENDENT SUBQUERY evaluation differs from **UNCACHEABLE SUBQUERY** evaluation. For **DEPENDENT SUBQUERY**, the subquery is re-evaluated only once for each set of different values of the variables from its outer context. For **UNCACHEABLE SUBQUERY**, the subquery is re-evaluated for each row of the outer context.

Cacheability of subqueries differs from caching of query results in the query cache (which is described in [Section 8.6.3.1. “How the Query Cache Operates”](#)). Subquery caching occurs during query execution, whereas the query cache is used to store results only after query execution finishes.

- **table**

The name of the table to which the row of output refers. This can also be one of the following values:

- **<unionM,N>**: The row refers to the union of the rows with **id** values of **M** and **N**.
- **<derivedN>**: The row refers to the derived table result for the row with an **id** value of **N**. A derived table may result, for example, from a subquery in the **FROM** clause.

- **type**

The join type. For descriptions of the different types, see [EXPLAIN Join Types](#).

- **possible_keys**

The **possible_keys** column indicates which indexes MySQL can choose from use to find the

rows in this table. Note that this column is totally independent of the order of the tables as displayed in the output from [EXPLAIN](#). That means that some of the keys in [possible_keys](#) might not be usable in practice with the generated table order.

If this column is [NULL](#), there are no relevant indexes. In this case, you may be able to improve the performance of your query by examining the [WHERE](#) clause to check whether it refers to some column or columns that would be suitable for indexing. If so, create an appropriate index and check the query with [EXPLAIN](#) again. See [Section 13.1.4, “ALTER TABLE Syntax”](#).

To see what indexes a table has, use [SHOW INDEX FROM tbl_name](#).

- [key](#)

The [key](#) column indicates the key (index) that MySQL actually decided to use. If MySQL decides to use one of the [possible_keys](#) indexes to look up rows, that index is listed as the key value.

It is possible that [key](#) will name an index that is not present in the [possible_keys](#) value. This can happen if none of the [possible_keys](#) indexes are suitable for looking up rows, but all the columns selected by the query are columns of some other index. That is, the named index covers the selected columns, so although it is not used to determine which rows to retrieve, an index scan is more efficient than a data row scan.

For [InnoDB](#), a secondary index might cover the selected columns even if the query also selects the primary key because [InnoDB](#) stores the primary key value with each secondary index. If [key](#) is [NULL](#), MySQL found no index to use for executing the query more efficiently.

To force MySQL to use or ignore an index listed in the [possible_keys](#) column, use [FORCE INDEX](#), [USE INDEX](#), or [IGNORE INDEX](#) in your query. See [Section 13.2.8.3, “Index Hint Syntax”](#).

For [MyISAM](#), [NDB](#), and [BDB](#) tables, running [ANALYZE TABLE](#) helps the optimizer choose better indexes. For [MyISAM](#) tables, [myisamchk --analyze](#) does the same as [ANALYZE TABLE](#). See [Section 7.6, “MyISAM Table Maintenance and Crash Recovery”](#).

- [key_len](#)

The [key_len](#) column indicates the length of the key that MySQL decided to use. The length is [NULL](#) if the [key](#) column says [NULL](#). Note that the value of [key_len](#) enables you to determine how many parts of a multiple-part key MySQL actually uses.

- [ref](#)

The [ref](#) column shows which columns or constants are compared to the index named in the [key](#) column to select rows from the table.

- [rows](#)

The [rows](#) column indicates the number of rows MySQL believes it must examine to execute the query.

For [InnoDB](#) tables, this number is an estimate, and may not always be exact.

- [Extra](#)

This column contains additional information about how MySQL resolves the query. For descriptions of the different values, see [EXPLAIN Extra Information](#).

EXPLAIN Join Types

The [type](#) column of [EXPLAIN](#) output describes how tables are joined. The following list describes the join

types, ordered from the best type to the worst:

- [system](#)

The table has only one row (= system table). This is a special case of the [const](#) join type.

- [const](#)

The table has at most one matching row, which is read at the start of the query. Because there is only one row, values from the column in this row can be regarded as constants by the rest of the optimizer. [const](#) tables are very fast because they are read only once.

[const](#) is used when you compare all parts of a [PRIMARY KEY](#) or [UNIQUE](#) index to constant values. In the following queries, *tbl_name* can be used as a [const](#) table:

```
SELECT * FROM tbl_name WHERE primary_key=1;
```

```
SELECT * FROM tbl_name
WHERE primary_key_part1=1 AND primary_key_part2=2;
```

- [eq_ref](#)

One row is read from this table for each combination of rows from the previous tables. Other than the [system](#) and [const](#) types, this is the best possible join type. It is used when all parts of an index are used by the join and the index is a [PRIMARY KEY](#) or [UNIQUE NOT NULL](#) index.

[eq_ref](#) can be used for indexed columns that are compared using the = operator. The comparison value can be a constant or an expression that uses columns from tables that are read before this table. In the following examples, MySQL can use an [eq_ref](#) join to process *ref_table*:

```
SELECT * FROM ref_table, other_table
WHERE ref_table.key_column=other_table.column;
```

```
SELECT * FROM ref_table, other_table
WHERE ref_table.key_column_part1=other_table.column
AND ref_table.key_column_part2=1;
```

- [ref](#)

All rows with matching index values are read from this table for each combination of rows from the previous tables. [ref](#) is used if the join uses only a leftmost prefix of the key or if the key is not a [PRIMARY KEY](#) or [UNIQUE](#) index (in other words, if the join cannot select a single row based on the key value). If the key that is used matches only a few rows, this is a good join type.

[ref](#) can be used for indexed columns that are compared using the = or <=> operator. In the following examples, MySQL can use a [ref](#) join to process *ref_table*:

```
SELECT * FROM ref_table WHERE key_column=expr;
```

```
SELECT * FROM ref_table, other_table
WHERE ref_table.key_column=other_table.column;
```

```
SELECT * FROM ref_table, other_table
WHERE ref_table.key_column_part1=other_table.column
AND ref_table.key_column_part2=1;
```

- [fulltext](#)

The join is performed using a **FULLTEXT** index.

- [ref_or_null](#)

This join type is like [ref](#), but with the addition that MySQL does an extra search for rows that contain **NULL** values. This join type optimization is used most often in resolving subqueries. In the following examples, MySQL can use a [ref_or_null](#) join to process *ref_table*:

```
SELECT * FROM ref_table
WHERE key_column=expr OR key_column IS NULL;
```

See [Section 8.3.1.6, “IS NULL Optimization”](#).

- [index_merge](#)

This join type indicates that the Index Merge optimization is used. In this case, the **key** column in the output row contains a list of indexes used, and **key_len** contains a list of the longest key parts for the indexes used. For more information, see [Section 8.3.1.4, “Index Merge Optimization”](#).

- [unique_subquery](#)

This type replaces [ref](#) for some **IN** subqueries of the following form:

```
value IN (SELECT primary_key FROM single_table WHERE some_expr)
```

[unique_subquery](#) is just an index lookup function that replaces the subquery completely for better efficiency.

- [index_subquery](#)

This join type is similar to [unique_subquery](#). It replaces **IN** subqueries, but it works for nonunique indexes in subqueries of the following form:

```
value IN (SELECT key_column FROM single_table WHERE some_expr)
```

- [range](#)

Only rows that are in a given range are retrieved, using an index to select the rows. The **key** column in the output row indicates which index is used. The **key_len** contains the longest key part that was used. The **ref** column is **NULL** for this type.

[range](#) can be used when a key column is compared to a constant using any of the **=**, **<>**, **>**, **>=**, **<**, **<=**, **IS NULL**, **<=>**, **BETWEEN**, or **IN()** operators:

```
SELECT * FROM tbl_name
WHERE key_column = 10;
```

```
SELECT * FROM tbl_name
WHERE key_column BETWEEN 10 and 20;
```

```
SELECT * FROM tbl_name
WHERE key_column IN (10,20,30);
```

```
SELECT * FROM tbl_name
```

```
WHERE key_part1= 10 AND key_part2 IN (10,20,30);
```

- [index](#)

The [index](#) join type is the same as [ALL](#), except that the index tree is scanned. This occurs two ways:

- If the index is a covering index for the queries and can be used to satisfy all data required from the table, only the index tree is scanned. In this case, the [Extra](#) column says [Using index](#). An index-only scan usually is faster than [ALL](#) because the size of the index usually is smaller than the table data.
- A full table scan is performed using reads from the index to look up data rows in index order. [Uses index](#) does not appear in the [Extra](#) column.

MySQL can use this join type when the query uses only columns that are part of a single index.

- [ALL](#)

A full table scan is done for each combination of rows from the previous tables. This is normally not good if the table is the first table not marked [const](#), and usually *very* bad in all other cases.

Normally, you can avoid [ALL](#) by adding indexes that enable row retrieval from the table based on constant values or column values from earlier tables.

EXPLAIN Extra Information

The [Extra](#) column of [EXPLAIN](#) output contains additional information about how MySQL resolves the query. The following list explains the values that can appear in this column. If you want to make your queries as fast as possible, look out for [Extra](#) values of [Using filesort](#) and [Using temporary](#).

- [const row not found](#)

For a query such as [SELECT ... FROM *tbl_name*](#), the table was empty.

- [Distinct](#)

MySQL is looking for distinct values, so it stops searching for more rows for the current row combination after it has found the first matching row.

- [Full scan on NULL key](#)

This occurs for subquery optimization as a fallback strategy when the optimizer cannot use an index-lookup access method.

- [Impossible HAVING](#)

The [HAVING](#) clause is always false and cannot select any rows.

- [Impossible WHERE](#)

The [WHERE](#) clause is always false and cannot select any rows.

- [Impossible WHERE noticed after reading const tables](#)

MySQL has read all [const](#) (and [system](#)) tables and notice that the [WHERE](#) clause is always false.

- [No matching min/max row](#)

No row satisfies the condition for a query such as [SELECT MIN\(...\) FROM ... WHERE *condition*](#).

- **no matching row in const table**

For a query with a join, there was an empty table or a table with no rows satisfying a unique index condition.

- **No tables used**

The query has no **FROM** clause, or has a **FROM DUAL** clause.

- **Not exists**

MySQL was able to do a **LEFT JOIN** optimization on the query and does not examine more rows in this table for the previous row combination after it finds one row that matches the **LEFT JOIN** criteria. Here is an example of the type of query that can be optimized this way:

```
SELECT * FROM t1 LEFT JOIN t2 ON t1.id=t2.id
WHERE t2.id IS NULL;
```

Assume that **t2.id** is defined as **NOT NULL**. In this case, MySQL scans **t1** and looks up the rows in **t2** using the values of **t1.id**. If MySQL finds a matching row in **t2**, it knows that **t2.id** can never be **NULL**, and does not scan through the rest of the rows in **t2** that have the same **id** value. In other words, for each row in **t1**, MySQL needs to do only a single lookup in **t2**, regardless of how many rows actually match in **t2**.

- **Range checked for each record (index map: N)**

MySQL found no good index to use, but found that some of indexes might be used after column values from preceding tables are known. For each row combination in the preceding tables, MySQL checks whether it is possible to use a **range** or **index merge** access method to retrieve rows. This is not very fast, but is faster than performing a join with no index at all. The applicability criteria are as described in [Section 8.3.1.3, “Range Optimization”](#), and [Section 8.3.1.4, “Index Merge Optimization”](#), with the exception that all column values for the preceding table are known and considered to be constants.

Indexes are numbered beginning with 1, in the same order as shown by **SHOW INDEX** for the table. The index map value **N** is a bitmask value that indicates which indexes are candidates. For example, a value of **0x19** (binary 11001) means that indexes 1, 4, and 5 will be considered.

- **Select tables optimized away**

The query contained only aggregate functions (**MIN()**, **MAX()**) that were all resolved using an index, or **COUNT(*)** for **MyISAM**, and no **GROUP BY** clause. The optimizer determined that only one row should be returned.

- **unique row not found**

For a query such as **SELECT ... FROM tbl_name**, no rows satisfy the condition for a **UNIQUE** index or **PRIMARY KEY** on the table.

- **Using filesort**

MySQL must do an extra pass to find out how to retrieve the rows in sorted order. The sort is done by going through all rows according to the join type and storing the sort key and pointer to the row for all rows that match the **WHERE** clause. The keys then are sorted and the rows are retrieved in sorted order. See [Section 8.3.1.11, “ORDER BY Optimization”](#).

- **Using index**

The column information is retrieved from the table using only information in the index tree without

having to do an additional seek to read the actual row. This strategy can be used when the query uses only columns that are part of a single index.

If the **Extra** column also says **Using where**, it means the index is being used to perform lookups of key values. Without **Using where**, the optimizer may be reading the index to avoid reading data rows but not using it for lookups. For example, if the index is a covering index for the query, the optimizer may scan it without using it for lookups.

- **Using index for group-by**

Similar to the **Using index** table access method, **Using index for group-by** indicates that MySQL found an index that can be used to retrieve all columns of a **GROUP BY** or **DISTINCT** query without any extra disk access to the actual table. Additionally, the index is used in the most efficient way so that for each group, only a few index entries are read. For details, see [Section 8.3.1.12, “GROUP BY Optimization”](#).

- **Using sort_union(...), Using union(...), Using intersect(...)**

These indicate how index scans are merged for the **index_merge** join type. See [Section 8.3.1.4, “Index Merge Optimization”](#).

- **Using temporary**

To resolve the query, MySQL needs to create a temporary table to hold the result. This typically happens if the query contains **GROUP BY** and **ORDER BY** clauses that list columns differently.

- **Using where**

A **WHERE** clause is used to restrict which rows to match against the next table or send to the client. Unless you specifically intend to fetch or examine all rows from the table, you may have something wrong in your query if the **Extra** value is not **Using where** and the table join type is **ALL** or **index**. Even if you are using an index for all parts of a **WHERE** clause, you may see **Using where** if the column can be **NULL**.

- **Using where with pushed condition**

This item applies to **NDBCLUSTER** tables *only*. It means that MySQL Cluster is using the Condition Pushdown optimization to improve the efficiency of a direct comparison between a nonindexed column and a constant. In such cases, the condition is “pushed down” to the cluster's data nodes and is evaluated on all data nodes simultaneously. This eliminates the need to send nonmatching rows over the network, and can speed up such queries by a factor of 5 to 10 times over cases where Condition Pushdown could be but is not used. For more information, see [Section 8.3.1.5, “Engine Condition Pushdown Optimization”](#).

EXPLAIN Output Interpretation

You can get a good indication of how good a join is by taking the product of the values in the **rows** column of the **EXPLAIN** output. This should tell you roughly how many rows MySQL must examine to execute the query. If you restrict queries with the **max_join_size** system variable, this row product also is used to determine which multiple-table **SELECT** statements to execute and which to abort. See [Section 8.9.2, “Tuning Server Parameters”](#).

The following example shows how a multiple-table join can be optimized progressively based on the information provided by **EXPLAIN**.

Suppose that you have the **SELECT** statement shown here and that you plan to examine it using **EXPLAIN**:

```
EXPLAIN SELECT tt.TicketNumber, tt.TimeIn,
```



```

tt.ProjectReference, tt.EstimatedShipDate,
tt.ActualShipDate, tt.ClientID,
tt.ServiceCodes, tt.RepetitiveID,
tt.CurrentProcess, tt.CurrentDPPerson,
tt.RecordVolume, tt.DPPrinted, et.COUNTRY,
et_1.COUNTRY, do.CUSTNAME
FROM tt, et, et AS et_1, do
WHERE tt.SubmitTime IS NULL
      AND tt.ActualPC = et.EMPLOYID
      AND tt.AssignedPC = et_1.EMPLOYID
      AND tt.ClientID = do.CUSTNMBR;

```

For this example, make the following assumptions:

- The columns being compared have been declared as follows.

| Table | Column | Data Type |
|-------|------------|-----------|
| tt | ActualPC | CHAR(10) |
| tt | AssignedPC | CHAR(10) |
| tt | ClientID | CHAR(10) |
| et | EMPLOYID | CHAR(15) |
| do | CUSTNMBR | CHAR(15) |

- The tables have the following indexes.

| Table | Index |
|-------|------------------------|
| tt | ActualPC |
| tt | AssignedPC |
| tt | ClientID |
| et | EMPLOYID (primary key) |
| do | CUSTNMBR (primary key) |

- The `tt.ActualPC` values are not evenly distributed.

Initially, before any optimizations have been performed, the [EXPLAIN](#) statement produces the following information:

```

table type possible_keys key  key_len ref  rows  Extra
et     ALL  PRIMARY          NULL NULL    NULL  74
do     ALL  PRIMARY          NULL NULL    NULL 2135
et_1   ALL  PRIMARY          NULL NULL    NULL  74
tt     ALL  AssignedPC,      NULL NULL    NULL 3872
      ClientID,
      ActualPC
Range checked for each record (index map: 0x23)

```

Because `type` is **ALL** for each table, this output indicates that MySQL is generating a Cartesian product of all the tables; that is, every combination of rows. This takes quite a long time, because the product of the number of rows in each table must be examined. For the case at hand, this product is $74 \times 2135 \times 74 \times 3872 = 45,268,558,720$ rows. If the tables were bigger, you can only imagine how long it would take.

One problem here is that MySQL can use indexes on columns more efficiently if they are declared as the

same type and size. In this context, [VARCHAR](#) and [CHAR](#) are considered the same if they are declared as the same size. `tt.ActualPC` is declared as `CHAR(10)` and `et.EMPLOYID` is `CHAR(15)`, so there is a length mismatch.

To fix this disparity between column lengths, use [ALTER TABLE](#) to lengthen `ActualPC` from 10 characters to 15 characters:

```
mysql> ALTER TABLE tt MODIFY ActualPC VARCHAR(15);
```

Now `tt.ActualPC` and `et.EMPLOYID` are both `VARCHAR(15)`. Executing the [EXPLAIN](#) statement again produces this result:

| table | type | possible_keys | key | key_len | ref | rows | Extra |
|-------|--------|--|---------|---------|-------------|------|----------------|
| tt | ALL | AssignedPC, ClientID, ActualPC | NULL | NULL | NULL | 3872 | Using where |
| do | ALL | PRIMARY | NULL | NULL | NULL | 2135 | |
| | | Range checked for each record (index map: 0x1) | | | | | |
| et_1 | ALL | PRIMARY | NULL | NULL | NULL | 74 | |
| | | Range checked for each record (index map: 0x1) | | | | | |
| et | eq_ref | PRIMARY | PRIMARY | 15 | tt.ActualPC | 1 | |

This is not perfect, but is much better: The product of the `rows` values is less by a factor of 74. This version executes in a couple of seconds.

A second alteration can be made to eliminate the column length mismatches for the `tt.AssignedPC = et_1.EMPLOYID` and `tt.ClientID = do.CUSTNMBR` comparisons:

```
mysql> ALTER TABLE tt MODIFY AssignedPC VARCHAR(15),
->          MODIFY ClientID VARCHAR(15);
```

After that modification, [EXPLAIN](#) produces the output shown here:

| table | type | possible_keys | key | key_len | ref | rows | Extra |
|-------|--------|--------------------------------------|----------|---------|---------------|------|----------------|
| et | ALL | PRIMARY | NULL | NULL | NULL | 74 | |
| tt | ref | AssignedPC, ClientID, ActualPC | ActualPC | 15 | et.EMPLOYID | 52 | Using where |
| et_1 | eq_ref | PRIMARY | PRIMARY | 15 | tt.AssignedPC | 1 | |
| do | eq_ref | PRIMARY | PRIMARY | 15 | tt.ClientID | 1 | |

At this point, the query is optimized almost as well as possible. The remaining problem is that, by default, MySQL assumes that values in the `tt.ActualPC` column are evenly distributed, and that is not the case for the `tt` table. Fortunately, it is easy to tell MySQL to analyze the key distribution:

```
mysql> ANALYZE TABLE tt;
```

With the additional index information, the join is perfect and [EXPLAIN](#) produces this result:

| table | type | possible_keys | key | key_len | ref | rows | Extra |
|-------|--------|-------------------------------------|---------|---------|-------------|------|----------------|
| tt | ALL | AssignedPC ClientID, ActualPC | NULL | NULL | NULL | 3872 | Using where |
| et | eq_ref | PRIMARY | PRIMARY | 15 | tt.ActualPC | 1 | |

```

et_1  eq_ref PRIMARY          PRIMARY 15          tt.AssignedPC 1
do    eq_ref PRIMARY          PRIMARY 15          tt.ClientID   1

```

Note that the **rows** column in the output from [EXPLAIN](#) is an educated guess from the MySQL join optimizer. You should check whether the numbers are even close to the truth by comparing the **rows** product with the actual number of rows that the query returns. If the numbers are quite different, you might get better performance by using [STRAIGHT_JOIN](#) in your [SELECT](#) statement and trying to list the tables in a different order in the [FROM](#) clause.

It is possible in some cases to execute statements that modify data when [EXPLAIN SELECT](#) is used with a subquery; for more information, see [Section 13.2.9.8, “Subqueries in the FROM Clause”](#).

[Previous](#) / [Next](#) / [Up](#) / [Table of Contents](#)

User Comments

Posted by Joost Boomkamp on February 8 2011 9:05am

[\[Delete\]](#) [\[Edit\]](#)

If you see the 'range checked for each record' message, it may be useful to verify that you're using the correct data types for the columns you are using in your query.

I got that message while I was searching for a way to speed up a slow table. then I noticed that one of the columns mentioned in the message had a VARCHAR datatype instead of INT (ouch)...

After fixing that, performance obviously much improved, and the range checked message went away.

Add your own comment

[Top](#) / [Previous](#) / [Next](#) / [Up](#) / [Table of Contents](#)



© 2014, Oracle Corporation and/or its affiliates